

# Conversational News Bot - Fullstack Prototyp

Bc.Phan Thi Anh Tuyet

# API Design for Real RAG

- Potential contract to call real RAG service:
  - Request:
    - query: user question
    - user\_id: for personalization & history
    - conversation\_id: to send previous context
  - Response:
    - answer: generated text
    - sources: list of { title, url, snippet }
    - latency\_ms, maybe model\_id
- Could define as OpenAPI/Swagger snippet for future integration

# Error Handling & UX

- Backend:
  - Simplicity: mocked data → fewer failure modes
- Frontend:
  - try/catch around fetch
  - On error: assistant message

**"Something went wrong when calling the API. Please try again."**
  - Loading state always cleared in finally
- UX:
  - User always sees *some* feedback (loading → answer or error)

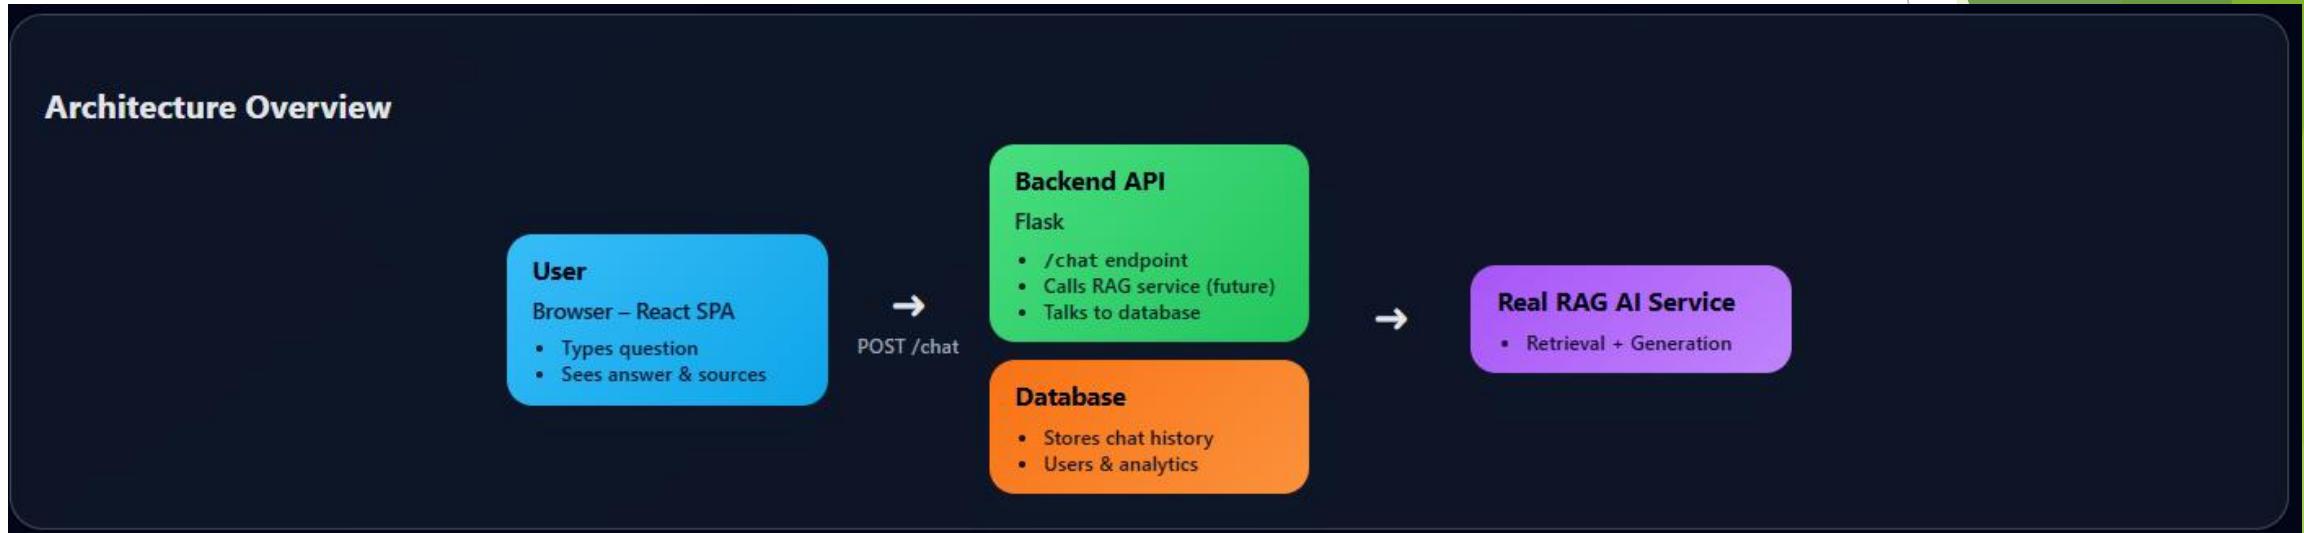
# State Management

- Options for storing chat history:
- **LocalStorage, Cookies, SessionStorage:**
  - **WOULD NOT RECOMMEND**, maybe for testing
    - Not shared between devices
- **Redis cache:**
  - Very fast
  - Good for ephemeral session data
  - Not ideal for long-term analytics
  - Would recommend for caching
- **Postgres (or similar DB):**
  - Persistent
  - Good for analytics, audit, per-user history
  - Slightly more complex schema & operations
  - Would recommend – MySQL, MariaDB, ORACLE DB, POSTGRE SQL

# Testing Approach

- Backend tests:
- Unit test for mocked RAG function:
  - Given question → returns deterministic answer + sources
- Integration test for /chat endpoint:
  - POST with a sample question → expect JSON with answer and sources
  - Explicit tests for "hello" and "thank u" flows
- Frontend tests (optional if by time):
- Component test:
  - Typing text and clicking Send adds user message
  - Mocked fetch → adds assistant message
- Snapshot/UI tests for chat bubbles

# High-Level Architecture



# Deployment & Scalability

- Production deployment idea:
- Package Flask app as Docker image
- Run in container platform (e.g. Kubernetes or simple VM)
- Frontend built as static bundle (React) and served via Nginx
- Most likely bottleneck:
  - Concurrent HTTP requests to the API
- Use:
  - Horizontal scaling (multiple replicas behind a load balancer)
  - Connection pooling to downstream services (RAG / DB)
- Real RAG latency:
  - Timeouts, retries, circuit breaker pattern at backend layer

# Live Demo

Thank you for your attention